

# LEARNING OBJECTIVES

- ◉ Define regularisation, bias and error metrics for regression problems
- ◉ Evaluate model fit using loss functions
- ◉ Select regression methods based on fit and complexity

---

## EVALUATING MODEL FIT

---

# PRE-WORK

---

## PRE-WORK REVIEW

---

- ◉ Understand goodness of fit (R-Squared)
- ◉ Measure statistical significance of features
- ◉ Recall what a residual is
- ◉ Implement a scikit-learn estimator to predict a target variable

---

**OPENING**

---

# **R-SQUARED AND RESIDUALS**

---

# OPENING

---



---

## WHAT IS R-SQUARED? WHAT IS A RESIDUAL?

---

- ◉ R-Squared is the central metric introduced for linear regression
- ◉ Which model performed better?
  - ◉ One with an R-Squared of 0.79 or 0.81?
- ◉ R-Squared measures explain variance
- ◉ But does it tell the magnitude or scale of error?
- ◉ We will explore loss functions and find ways to refine our model


## WHAT IS R-SQUARED? WHAT IS A RESIDUAL?

- ◉ In a data set with  $n$  cases where the outputs are  $y_i = y_1, \dots, y_n$
- ◉ And the predicted values are (note the  $y$ -hat)  $\hat{y}_i = \hat{y}_1, \dots, \hat{y}_n$
- ◉ The mean of  $y$  is (note the  $y$ -bar)

$$\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$$

$$SS_{tot} = SS_{reg} + SS_{res}$$

$$R^2 = \frac{SS_{reg}}{SS_{tot}}$$




# WHAT IS R-SQUARED? WHAT IS A RESIDUAL?

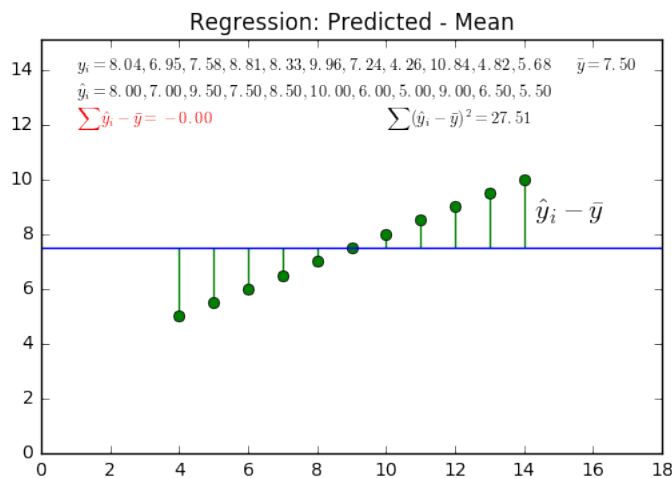
## ◉ The Sum of Squares

**Total SS** = **Regression SS** + **Residual SS**


$$SS_{tot} = \sum_{i=1}^n (y_i - \bar{y})^2$$



$$SS_{reg} = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2$$



$$SS_{res} = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$





---


**INTRODUCTION**

---

# **LINEAR MODELS AND ERROR**

## RECALL: WHAT IS RESIDUAL ERROR?

- ◉ In linear models, residual error must be normal with a median close to zero
- ◉ Individual residuals are useful to see the error of specific points, but it does not provide an overall picture for optimisation
- ◉ We need a metric to summarise the error in our model into one value
  - ◉ **Mean Square Error**: the mean residual error in our model



---

## MEAN SQUARE ERROR (MSE)

---

- ◉ To calculate MSE
  - ◉ Calculate the difference between each target  $y$  and the model's predicted value  $\hat{y}$  (i.e. the residual)
  - ◉ Square each residual
  - ◉ Take the mean of the squared residual errors

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

---

## MEAN SQUARE ERROR (MSE)

---

- ◉ scikit-learn's metrics module includes a `mean_squared_error()` function
- ◉ For example, two arrays of the same values would have an MSE of 0
- ◉ Two arrays with different values would have a positive MSE

```
from sklearn import metrics
# metrics.mean_squared_error(y, model.predict(X))

print(metrics.mean_squared_error([1, 2, 3, 4, 5], [1, 2, 3, 4, 5]))
0.0

print(metrics.mean_squared_error([1, 2, 3, 4, 5], [5, 4, 3, 2, 1]))
# (4**2 + 2**2 + 0**2 + 2**2 + 4**2) / 5
8.0
```

---

## HOW DO WE MINIMISE THE ERROR?

---

- ◉ The regression method we have used is called
  - ◉ “Ordinary Least Squares”
- ◉ This means that given a matrix  $X$ 
  - ◉ Solve for the least amount of square error for  $y$
- ◉ However, this assumes that  $X$  is unbiased
  - ◉ That it is representative of the population

## LET'S COMPARE TWO RANDOM MODELS

```
import numpy as np
import pandas as pd
from sklearn import linear_model, metrics

np.random.seed(seed = 100)
df = pd.DataFrame({"x": range(100), "y": range(100)})
biased_df = df.copy()
biased_df.loc[:20, "x"], biased_df.loc[:20, "y"] = 1, 1

def append_jitter(series):
    return series + np.random.random_sample(size = 100)

df["x"], df["y"] = append_jitter(df.x), append_jitter(df.y)
biased_df["x"], biased_df["y"] = append_jitter(biased_df.x), append_jitter(biased_df.y)


# Fit:
lm = linear_model.LinearRegression().fit(df[["x"]], df["y"])
print(metrics.mean_squared_error(df["y"], lm.predict(df[["x"]])))

# Biased fit:
lm = linear_model.LinearRegression().fit(biased_df[["x"]], biased_df["y"])
print(metrics.mean_squared_error(df["y"], lm.predict(df[["x"]])))
```

0.175065131142

0.187115987374

# BIAS VS VARIANCE



---


## BIAS / VARIANCE TRADEOFF

---

- ◉ When our error is **biased**, it means the model's prediction is consistently far away from the actual value
  - ◉ This could be a sign of poor sampling and poor data
- ◉ One objective of a biased model is to trade bias error for generalised error
  - ◉ We prefer the error to be more evenly distributed across the model
- ◉ This is called error due to **variance**
- ◉ We want our model to generalise to data it has not seen even if does not perform as well on data it has already seen



# BIAS / VARIANCE TRADEOFF



---

## ACTIVITY: KNOWLEDGE CHECK

---

**DIRECTIONS: ANSWER THE FOLLOWING QUESTIONS (5 MINUTES)**

---

1. Which of the following scenarios would be better for a weather reporter?
  - a. Knowing that I can very accurately “predict” the temperature outside from previous days perfectly, but be 20-30 degrees off for future days
  - b. Knowing that I can accurately “predict” the general trend of the temperature outside from previous days and therefore am at most only 10 degrees off on future days



EXERCISE

---

**DEMONSTRATION**

---

# CROSS VALIDATION

---

## CROSS VALIDATION

---

- ◉ Cross Validation can help account for bias
- ◉ The general idea is to
  - ◉ Generate several models on different cross sections of the data
  - ◉ Measure the performance of each
  - ◉ Take the mean performance
- ◉ This technique swaps bias error for generalised error, describing previous trends accurately enough to extend to future trends

---

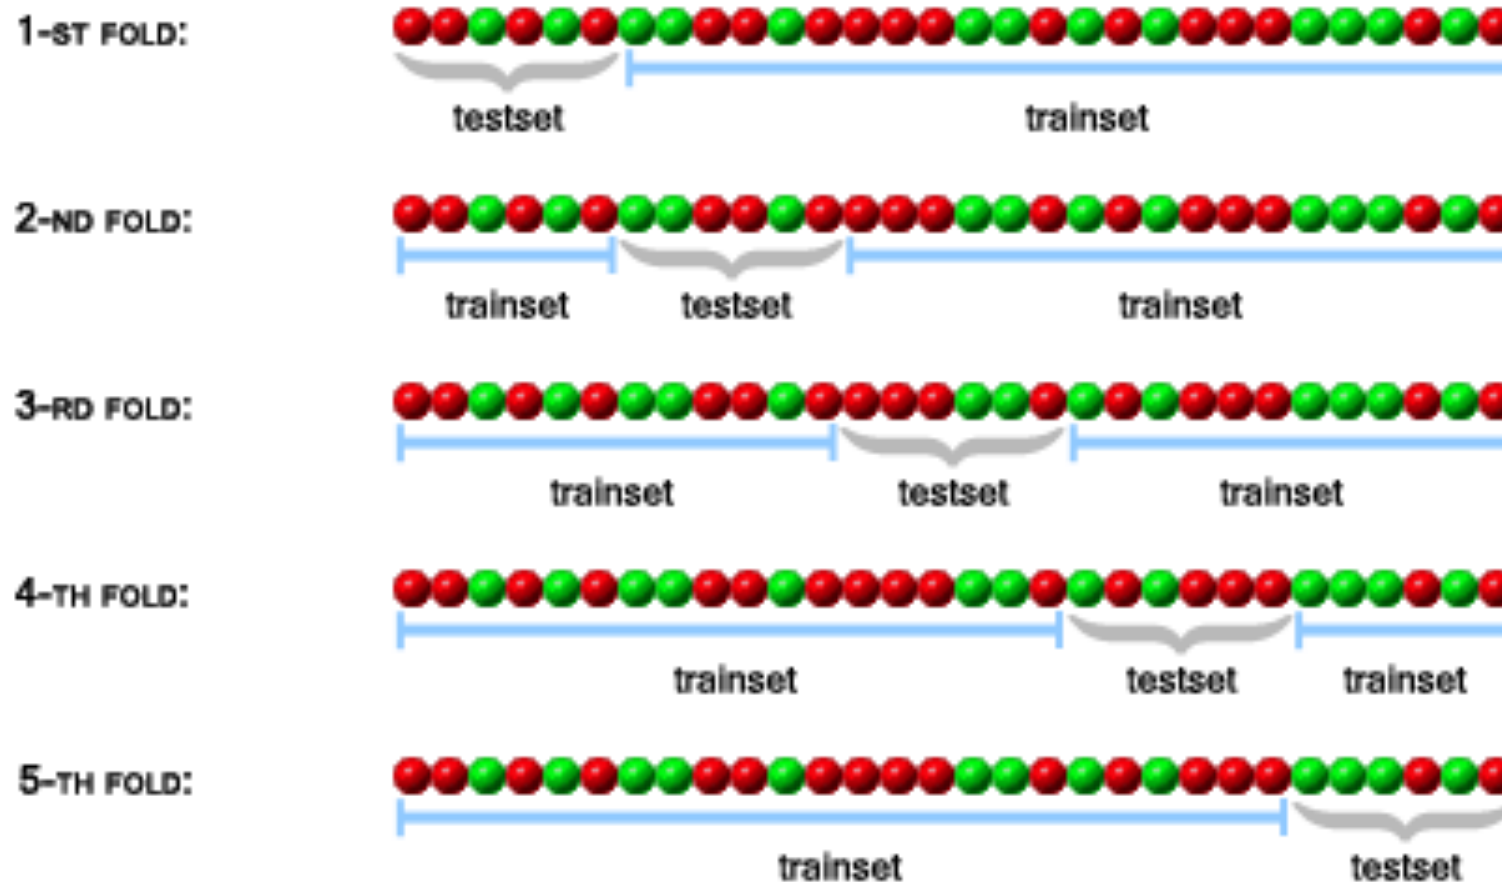
## K-FOLD CROSS VALIDATION

---

- ◉ k-fold cross validation
  - ◉ Split the data into k group
  - ◉ Train the model on all segments except one
  - ◉ Test model performance on the remaining set
- ◉ If  $k = 5$ , split the data into five segments and generate five models

# CROSS VALIDATION

## ONE ITERATION OF A 5-FOLD CROSS-VALIDATION:



## USING K-FOLD CROSS VALIDATION WITH MSE

- ## • Import the appropriate packages and load data

```
from sklearn import cross_validation, metrics
np.random.seed(seed = 100)

wd      = "../assets/dataset/"
bikeshare = pd.read_csv(wd + "bikeshare.csv")
weather  = pd.get_dummies(bikeshare.weathersit, prefix = "weather")
modeldata = bikeshare[["temp", "hum"]].join(weather[["weather_1",
                                                    "weather_2",
                                                    "weather_3"]])
y        = bikeshare.casual
```

---

## USING K-FOLD CROSS VALIDATION WITH MSE

---

- © Build models on subsets of the data and calculate the average score

```
kf = cross_validation.KFold(len(modeldata), n_folds = 5, shuffle = True)

scores = []
for train_index, test_index in kf:
    lm = linear_model.LinearRegression().fit(
        modeldata.iloc[train_index],
        y.iloc[train_index])
    scores.append(metrics.mean_squared_error(
        y.iloc[test_index],
        lm.predict(modeldata.iloc[test_index])))
```



---

## USING K-FOLD CROSS VALIDATION WITH MSE

---

- ◉ This can be compared to the model built on all of the data

```
print(scores)
print(np.mean(scores))
# This score will be lower
# but we are trading off bias error for generalised error:
lm = linear_model.LinearRegression().fit(modeldata, y)
print(metrics.mean_squared_error(y, lm.predict(modeldata)))
```

```
[1432.81283012, 1758.18152140, 1706.87054437, 1808.52000170, 1662.48998234]
1673.77497599
1672.58110765
```

- ◉ Which approach would predict new data more accurately?

---

**GUIDED PRACTICE**

---

# **CROSS VALIDATION WITH LINEAR REGRESSION**

---

## ACTIVITY: KNOWLEDGE CHECK

---



### EXERCISE

#### **DIRECTIONS: (20 MINUTE)**

---

If we were to continue increasing the number of folds in cross validation, would error increase or decrease?

1. Using the previous code example, perform k-fold cross validation for all even numbers between 2 and 50.
2. Answer the following questions:
  - a. What does `shuffle = True` do?
  - b. At what point does cross validation no longer seem to help the model?
3. Hint: `range(2, 51, 2)` produces a list of even numbers from 2 to 50

---

**INTRODUCTION**

---

**REGULARISATION**

**AND**

**CROSS VALIDATION**

---


## WHAT IS REGULARISATION AND WHY USE IT?

---

- ◉ Regularisation is an additive approach to protect models against overfitting (being potentially biased and overconfident, not generalising well)
- ◉ Regularisation becomes an additional weight to coefficients, shrinking them closer to zero
- ◉ L1 (Lasso Regression) adds the extra weight to coefficients
  - ◉ address primarily reducing features by making coefficients zero
- ◉ L2 (Ridge Regression) adds the square of the extra weight to coefficients
  - ◉ address primarily reducing outliers' influence

# WHAT IS OVERFITTING?

- ◉ The first model poorly explains the data
- ◉ The second model explains the general curve of the data
- ◉ The third model drastically overfits the model, bending to every point



- ◉ Regularisation helps prevent the third model

## WHERE REGULARISATION MAKES SENSE

- ◉ What happens to MSE if use Lasso or Ridge Regression directly?

```
lm = linear_model.LinearRegression().fit(modeldata, y)
print(metrics.mean_squared_error(y, lm.predict(modeldata)))
lm = linear_model.Lasso().fit(modeldata, y)
print(metrics.mean_squared_error(y, lm.predict(modeldata)))
lm = linear_model.Ridge().fit(modeldata, y)
print(metrics.mean_squared_error(y, lm.predict(modeldata)))
```

1672.58110765	OLS
1725.41581608	L1
1672.60490113	L2

---

## WHERE REGULARISATION MAKES SENSE

---

- ◉ It does not seem to help
  - ◉ Why is that?
- ◉ We need to optimise the regularisation weight parameter (called  $\alpha$ ) through cross validation



---

## ACTIVITY: KNOWLEDGE CHECK

---

**DIRECTIONS: ANSWER THE FOLLOWING QUESTIONS (5 MINUTES)**

---

1. Why is regularisation important?
2. What does it protect against and how?



EXERCISE

---

**DEMONSTRATION**

---

# **UNDERSTANDING REGULARISATION EFFECTS**

---

## QUICK CHECK

---

- ◉ We are working with the bikeshare data to predict riders over hours/days with a few features
- ◉ Does it make sense to use a ridge regression or a lasso regression?
- ◉ Why?

---

## UNDERSTANDING REGULARISATION EFFECTS

---

- ◉ Let's test a variety of alpha weights for Ridge Regression on the bikeshare data

```
alphas = np.logspace(-10, 10, 21)
for a in alphas:
    print("Alpha:", a)
    lm = linear_model.Ridge(alpha = a)
    lm.fit(modeldata, y)
    print(lm.coef_)
    print(metrics.mean_squared_error(y, lm.predict(modeldata)))
```

- ◉ What happens to the weights of the coefficients as alpha increases?
- ◉ What happens to the error as alpha increases?

---

## EASIER WITH GRID SEARCH

---

- ◉ Grid search exhaustively searches through all given options to find the best solution
- ◉ Grid search will try all combos given in param\_grid

```
param_grid = {  
    "intercept": [True, False],  
    "alpha": [1, 2, 3]  
}
```

## EASIER WITH GRID SEARCH

```
param_grid = {  
    "intercept": [True, False],  
    "alpha": [1, 2, 3]  
}
```

intercept	alpha
True	1
True	2
True	3
False	1
False	2
False	3

---

## EASIER WITH GRID SEARCH

---

- © This is an incredibly powerful, automated machine learning tool!

```
from sklearn import grid_search

alphas = np.logspace(-10, 10, 21)
gs = grid_search.GridSearchCV(
    estimator = linear_model.Ridge(),
    param_grid = {"alpha": alphas},
    scoring = "mean_squared_error")
```

---

## EASIER WITH GRID SEARCH

---

- © This is an incredibly powerful, automated machine learning tool!

```
gs.fit(modeldata, y)

# mean squared error here comes in negative, so let's make it positive.
print(-gs.best_score_)

# explains which grid_search setup worked best
print(gs.best_estimator_)

# shows all the grid pairings and their performances
print(gs.grid_scores_)
```



---

## GUIDED PRACTICE

---

**GRID SEARCH  
CROSS VALIDATION,  
SOLVING FOR ALPHA**

---

## ACTIVITY: GRID SEARCH CROSS VALIDATION, SOLVING FOR ALPHA

---



### EXERCISE

**DIRECTIONS: ANSWER THE FOLLOWING QUESTIONS (25 MINUTES)**

---

1. Modify the previous code to do the following:
  - a. Introduce cross validation into the grid search
    - i. This is accessible from the `cv` argument
  - b. Add `fit_intercept = True` and `False` to the `param_grid` dictionary
  - c. Re-investigate the best score, best estimator and grid score attributes as a result of the grid search

## INTRODUCTION

---

# MINIMISING LOSS THROUGH GRADIENT DESCENT


---

## GRADIENT DESCENT


---

- ◉ Gradient Descent can also help us minimise error
- ◉ How Gradient Descent works
  - ◉ A random linear solution is provided as a starting point
  - ◉ The solver attempts to find a next “step”: take a step in any direction and measure the performance
  - ◉ If the solver finds a better solution (i.e. lower MSE), this is the new starting point
  - ◉ Repeat these steps until the performance is optimised and no “next steps” perform better
    - ◉ The size of steps will shrink over time


# GRADIENT DESCENT



# GRADIENT DESCENT



# GRADIENT DESCENT



## A CODE EXAMPLE OF GRADIENT DESCENT

```
num_to_approach, start, steps, optimised = 6.2, 0., [-1, 1], False
while not optimised:
    current_distance = num_to_approach - start
    got_better = False
    next_steps = [start + i for i in steps]
    for n in next_steps:
        distance = np.abs(num_to_approach - n)
        if distance < current_distance:
            got_better = True
            print(distance, "is better than", current_distance)
            current_distance = distance
            start = n
    if got_better:
        print("found better solution! using ", current_distance)
        a += 1
    else:
        optimised = True
        print(start, " is closest to ", num_to_approach)
```



---


## A CODE EXAMPLE OF GRADIENT DESCENT

---

- What is the code doing?
- What could go wrong?

# GLOBAL VS LOCAL MINIMUMS

- ◉ Gradient Descent could solve for a **local** minimum instead of a **global** minimum
- ◉ A local minimum is confined to a very specific subset of solutions
- ◉ The global minimum considers all solutions
- ◉ These could be equal, but that is not always true



---

**DEMONSTRATION**

---

# APPLICATION OF GRADIENT DESCENT

---

## APPLICATION OF GRADIENT DESCENT

---

- ◉ Gradient Descent works best when
  - ◉ We are working with a large dataset
    - ◉ Smaller datasets are more prone to error
  - ◉ Data is cleaned up and normalised
- ◉ Gradient Descent is significantly faster than OLS
  - ◉ This becomes important as data gets bigger

---

## APPLICATION OF GRADIENT DESCENT

---

- ◉ We can easily run a Gradient Descent regression
  - ◉ Note: The verbose argument can be set to 1 to see the optimisation steps.

```
lm = linear_model.SGDRegressor()  
lm.fit(modeldata, y)  
print(lm.score(modeldata, y))  
print(metrics.mean_squared_error(y, lm.predict(modeldata)))
```

- ◉ Untuned, how well did gradient descent perform compared to OLS?

---

## APPLICATION OF GRADIENT DESCENT

---

- ⦿ Gradient Descent can be tuned with
  - ⦿ the learning rate: how aggressively we solve the problem
  - ⦿ epsilon: at what point do we say the error margin is acceptable
  - ⦿ iterations: when should we stop no matter what

---

**INDEPENDENT PRACTICE**

---

**HANDS ON**

# ACTIVITY: HANDS ON

---



## EXERCISE

### DIRECTIONS (20 MINUTES)

---

There are tons of ways to approach a regression problem.

1. Implement the Gradient Descent approach to our bikeshare modelling problem
2. Show how Gradient Descent solves and optimises the solution
3. Demonstrate the `grid_search` module
4. Use a model you evaluated last class or the simpler one from today.
5. Implement `param_grid` in grid search to answer the following questions:
  - a. With a set of values between  $10^{-10}$  and  $10^{-1}$ , how does MSE change?
  - b. Our data suggests we use L1 regularisation. Using a grid search with `l1_ratios` between 0 and 1, increasing every 0.05, does this statement hold true? If not, did gradient descent have enough iterations to work properly?



## ACTIVITY: HANDS ON



### EXERCISE

```
params = {} # put your gradient descent parameters here

gs = grid_search.GridSearchCV(
    estimator = linear_model.SGDRegressor(),
    cv        = cross_validation.KFold(len(modeldata),
    n_folds    = 5,
    shuffle    = True),
    param_grid = params,
    scoring    = "mean_squared_error")

gs.fit(modeldata, y)

print("BEST ESTIMATORS")
print(-gs.best_score_)
print(gs.best_estimator_)
print("ALL ESTIMATORS")
print(gs.grid_scores_)
```

---

**CONCLUSION**

---

# TOPIC REVIEW

---

## TOPIC REVIEW

---

- ◉ What is the (typical) range of R-Squared?
- ◉ What is the range of Mean Squared Error (MSE)?
- ◉ How would changing the scale or interpretation of  $y$  (your target variable) effect the Mean Squared Error?
- ◉ What is Cross Validation and why do we use it in Machine Learning?
- ◉ What is error due to Bias?
- ◉ What is error due to Variance?
- ◉ Which is better for a model to have, if it had to have one?
- ◉ How does Gradient Descent try a different approach to minimising error?

**DATA SCIENCE**

---

**BEFORE NEXT CLASS**

---

## **BEFORE NEXT CLASS**

---

# **DUE DATE**

- ◉ Project
  - ◉ Final Project, part 1

---

## EVALUATING MODEL FIT


---

Q & A

# **CREDITS AND REFERENCES**

# EVALUATING MODEL FIT

---



- “[In Science] All models are wrong; some are useful!”
  - [George Box](#), Wikipedia



- All models are wrong: limits of science
  - [Martin Hilbert](#), University of California, Davis
  - [youTube](#), video