

LATENT VARIABLES AND NLP

Angelo Klin
Katra Analytics

LEARNING OBJECTIVES

- ◉ Understand what **latent variables** are
- ◉ Understand the uses of latent variables in language processing
- ◉ Use the **word2vec** and **LDA** algorithms of **gensim**

LATENT VARIABLES AND NLP

PRE-WORK

PRE-WORK REVIEW

- ◉ Install [gensim](#) with

```
conda install gensim
```

- ◉ Recall and apply unsupervised learning techniques
- ◉ Recall probability distributions, specifically discrete multinomial distributions
- ◉ Recall NLP essentials, including experience with [spaCy](#)
- ◉ **BONUS:** Setup Twitter API credentials using the provided instructions

OPENING

LATENT VARIABLE MODELS

LATENT VARIABLE MODELS

- ◎ This lesson will continue on natural language processing with an emphasis on latent variables models
- ◎ Mining and Refining data is a key part of the data science workflow
- ◎ In our last class, we saw many techniques for mining the data, including preprocessing, building linguistic rules to uncover patterns and creating classifiers from unstructured data
- ◎ In this class, we will continue with methods to Refine our understanding of the text by attempting to uncover structure or organisation in the text

LATENT VARIABLE MODELS


- ◉ Many advances in NLP are based on using data to learn rules of grammar and language
- ◉ We saw these tools in our last class
 - ◉ Tokenisation
 - ◉ Stemming or lemmatisation
 - ◉ Parsing and tagging
- ◉ Each of these are based on a classical or theoretical understanding of language

LATENT VARIABLE MODELS

- Tokenisation:

- John hit the ball

- [John, hit, the, ball]



- Where did you go

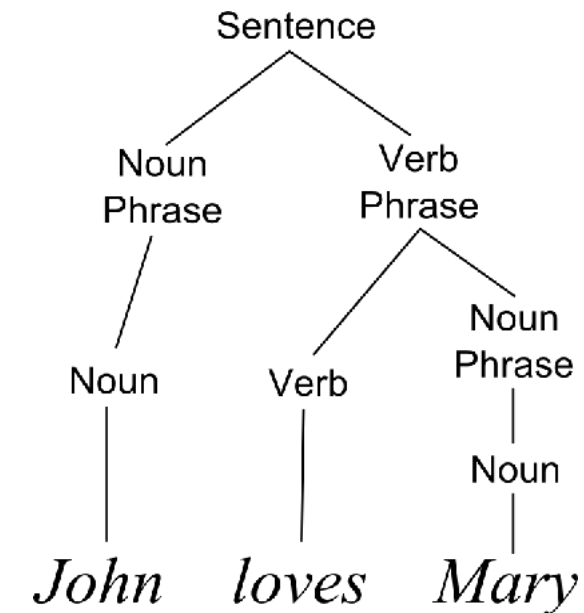
- [Where, did, you, go]

- Stemming or lemmatisation:

- shouted → shout

- better → good

- Parsing and tagging:



LATENT VARIABLE MODELS

- ◉ Latent variable models are different in that they try to understand language based on how the words are used
- ◉ For example, instead of learning that “bad” and “badly” are related because they share the same root, we will determine that they are related because they are often used in the same way often or near the same words
- ◉ We will use unsupervised techniques (discovering patterns or structure) to extract the information

LATENT VARIABLE MODELS

Traditional NLP Models	Latent Variable Models
Focused on theoretical understanding of language	Focused on how the language is actually used in practice
Tries to learn the rules of a particular language	Infers meaning from how words are used together
Preprogrammed set of rules	Uses unsupervised learning to discover patterns or structure

LATENT VARIABLE MODELS

Traditional NLP Models	Latent Variable Models
“bad” and “badly” are related because they share a common root	“bad” and “badly” are related because they are used the same way or near the same words
“Python” and “C++” are both programming languages because they are often a noun preceded by the verb “program” or “code	“Python” and “C++” are both programming languages because they are often used in the same context

INTRODUCTION

LATENT VARIABLE MODELS

LATENT VARIABLE MODELS

- ◉ Latent variable models are models in which we assume the data we are observing has some hidden, underlying structure that we can not see and which we would like to learn
- ◉ These hidden, underlying structures are the latent (i.e. hidden) variables we want our model to understand
- ◉ Text processing is a common application of latent variables


LATENT VARIABLE MODELS

- ◉ While language (in the classical sense) is defined by a set of pre-structured grammar rules and a vocabulary, we often break those rules and create new words (e.g. selfie)
- ◉ Instead of attempting to train our model on the rules of proper grammar, we will ignore grammar and seek to uncover alternate hidden structures

LATENT VARIABLE MODELS






- © Latent variable techniques are often used for recommending news articles or mining large troves of data to find commonalities
- © Topic modelling, a method we will cover today, is used in [the NY times recommendation engine](#)
- © The New York Times attempts to map their articles to a latent space of topics using the content of the article

LATENT VARIABLE MODELS



LATENT VARIABLE MODELS

- © [Lyst](#), an online fashion retailer, uses latent representations of clothing descriptions to find similar clothing

	Term	Similarity
	"riding"	" 0.962305
	"ankle"	0.888326
	"chelsea"	0.885039
	"gloss"	0.862084
	"combat"	0.861708

DIMENSIONALITY REDUCTION IN TEXT REPRESENTATION

- Our previous representation of a set of text documents (articles) for classification was a matrix with one row per document and one column per word (or n-gram)

The diagram illustrates the process of dimensionality reduction using matrix multiplication. It shows three matrices and their relationship:

- Document Matrix (MD):** A matrix with 7 rows and 5 columns. The columns are labeled "data", "inf.", "brain", "lung", and "retrieval". The rows are labeled "CS" (Concept Strengths) and "MD" (Matrix Dimensions).
- Doc-to-Concept Similarity Matrix:** A 7x2 matrix with values: $\begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix}$.
- Concepts Strengths:** A 2x2 matrix with values: $\begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix}$.
- Term-to-Concept Similarity Matrix:** A 5x6 matrix with values: $\begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$.

The equation shown is:

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix} \times \begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

Arrows point from the labels "doc-to-concept similarity matrix", "concepts strengths", and "term-to-concept similarity matrix" to their respective matrices. The "CS" and "MD" labels are placed next to the first matrix with arrows indicating the dimensions.

DIMENSIONALITY REDUCTION IN TEXT REPRESENTATION

- © While this sums up most of the information, it does drop a few things, mostly structure and order
- © Additionally, many of the columns may be correlated

DIMENSIONALITY REDUCTION IN TEXT REPRESENTATION

- © For example, an article that contains the word “IPO” is likely to contain the word “stock” or “ASX”
- © Therefore, those columns are repetitive and likely to represent the same concept or idea
- © For classification, we may only care that there are finance-related words

DIMENSIONALITY REDUCTION IN TEXT REPRESENTATION

- ◉ One way to deal with this is through regularisation - L1/Lasso regularisation tends to remove repetitive features by bringing their learned coefficients to 0
- ◉ Another is to perform dimensionality reduction, where we first identify the correlated columns and then replace them with a column that represents the concept they have in common
- ◉ For instance, we could replace “IPO”, “stocks” and “ASX” with a single column “HasFinancialWords”

DIMENSIONALITY REDUCTION IN TEXT REPRESENTATION

- ◉ There are many techniques to do this automatically and most follow a very similar approach
- ◉ Identify correlated columns
- ◉ Replace them with a new column that encapsulates the others

Doc #	Car	Truck	Van	Dog
6344	1	1	1	0
6345	0	1	1	1
6346	1	1	1	0



Doc #	Vehicle	Dog
6344	1	0
6345	1	1
6346	1	0

DIMENSIONALITY REDUCTION IN TEXT REPRESENTATION


- ◎ The techniques vary in how they define correlation and how much of the relationship between the original and new columns you need to save
- ◎ Dimensionality techniques can vary between linear and non-linear

DIMENSIONALITY REDUCTION IN TEXT REPRESENTATION

- ◉ There are many techniques build into scikit-learn
- ◉ One of the most common is **Principal Component Analysis** (PCA)
- ◉ PCA, when applied to text data, is sometimes known as **Latent Semantic Indexing** (LSI)

DIMENSIONALITY REDUCTION IN TEXT REPRESENTATION

- PCA helps reduce the feature space into fewer dimensions



MIXTURE MODELS AND LANGUAGE PROCESSING

- ◎ Mixture models (specifically **Latent Dirichlet Allocation** or LDA) take this concept further and generate more structure around the documents
- ◎ Instead of just replacing correlated columns, we create clusters of common words and generate probability distributions to explicitly state how related words are

MIXTURE MODELS AND LANGUAGE PROCESSING

- ◉ To understand this better, let's imagine a new way to generate text:
 - ◉ Start writing a document
 - ◉ Choose a topic (sports, news, science)
 - ◉ Choose a random word from that topic
 - ◉ Repeat
- ◉ Repeat for the next document

MIXTURE MODELS AND LANGUAGE PROCESSING

- ◉ This “model” of text is assuming that each document is some mixture of topics
- ◉ It may be mostly science but may contain some business information
- ◉ The latent structure we want to uncover are the topics (or concepts) that generate that text

MIXTURE MODELS AND LANGUAGE PROCESSING

- ◉ Latent Dirichlet Allocation is a model that assumes this is the way text is generated and then attempts to learn two things:
 - ◉ The word distribution of each topic
 - ◉ The topic distribution of each document

MIXTURE MODELS AND LANGUAGE PROCESSING

Topics

gene	0.04
dna	0.02
genetic	0.01
...	

life	0.02
evolve	0.01
organism	0.01
...	

brain	0.04
neuron	0.02
nerve	0.01
...	

data	0.02
number	0.02
computer	0.01
...	

Documents

Seeking Life's Bare (Genetic) Necessities

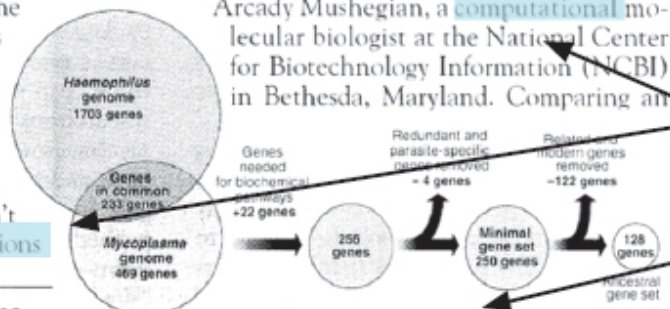
COLD SPRING HARBOR, NEW YORK—How many **genes** does an **organism** need to **survive**? Last week at the genome meeting here,* two genome researchers with radically different approaches presented complementary views of the basic genes needed for **life**. One research team, using **computer** analyses to compare known **genomes**, concluded that today's **organisms** can be sustained with just 250 genes, and that the earliest life forms required a mere 128 **genes**. The other researcher mapped genes in a simple parasite and estimated that for this organism, 800 genes are plenty to do the job—but that anything short of 100 wouldn't be enough.

Although the numbers don't match precisely, those **predictions**

"are not all that far apart," especially in comparison to the 75,000 **genes** in the human genome, notes Siv Andersson of Uppsala University in Sweden, who arrived at the 800 number. But coming up with a consensus answer may be more than just a **genetic numbers** game, particularly as more and more **genomes** are completely mapped and sequenced. "It may be a way of organizing any newly **sequenced genome**," explains Arcady Mushegian, a **computational** molecular biologist at the National Center for Biotechnology Information (NCBI) in Bethesda, Maryland. Comparing an


* Genome Mapping and Sequencing, Cold Spring Harbor, New York, May 8 to 12.

SCIENCE • VOL. 272 • 24 MAY 1996




Stripping down. **Computer analysis** yields an estimate of the minimum modern and ancient genomes.

Topic proportions and assignments



MIXTURE MODELS AND LANGUAGE PROCESSING

- The word distribution is a multinomial distribution of each topic representing what words are most likely from that topic



MIXTURE MODELS AND LANGUAGE PROCESSING

- ◉ For example, let's say we have three topics: sports, business and science
- ◉ For each topic, we uncover the most likely words to come from them:

```
sports: [football: 0.3, basketball: 0.2, baseball: 0.2, touchdown: 0.02 ... genetics: 0.0001]  
science: [genetics: 0.2, drug: 0.2, ... baseball: 0.0001]  
business: [stocks: 0.1, ipo: 0.08, ... baseball: 0.0001]
```

- ◉ For each word and topic pair, we learn some probability:
 - ◉ $P(\text{word} \mid \text{topic})$


MIXTURE MODELS AND LANGUAGE PROCESSING

- ◉ The topic distribution is a multinomial distribution for each document representing what topics are most likely to appear in that document
- ◉ For all our of sample documents, we have a distribution over {sports, science, business}:

```
ESPN article: [sports: 0.8, business: 0.2, science: 0.0]  
Bloomberg article: [business: 0.7, science: 0.2, sports: 0.1]
```

- ◉ For each topic and document pair, we learn some probability:
 - ◉ $P(\text{topic} \mid \text{document})$

MIXTURE MODELS AND LANGUAGE PROCESSING



“Genetics”

human
genome
dna
genetic
genes
sequence
gene
molecular
sequencing
map
information
genetics
mapping
project
sequences

“Evolution”

evolution
evolutionary
species
organisms
life
origin
biology
groups
phylogenetic
living
diversity
group
new
two
common

“Disease”

disease
host
bacteria
diseases
resistance
bacterial
new
strains
control
infectious
malaria
parasite
parasites
united
tuberculosis

“Computers”

computer
models
information
data
computers
system
network
systems
model
parallel
methods
networks
software
new
simulations

MIXTURE MODELS AND LANGUAGE PROCESSING

- ◉ Topic models are useful for organising a collection of documents and uncovering the main underlying concepts
- ◉ There are many variants that attempt to add even more structure to the “model”:
 - ◉ Supervised topic models guide the process with pre-decided topics
 - ◉ Position-dependent topic models ignore which words occur in which document and instead focus on where they occur
 - ◉ Variable number topic models test different numbers of topics to find the best model

ACTIVITY: KNOWLEDGE CHECK

DIRECTIONS: ANSWER THE FOLLOWING QUESTIONS

1. Take any recent news article and brainstorm which three topics this story is most likely to be made up of
2. Next, brainstorm which words are most likely derived from which of those three topics



EXERCISE

DEMONSTRATION

LDA IN GENSIM

LDA IN GENSIM

- © [gensim](#) is a library of language processing tools focused on latent variable models of text
- © It was originally developed by graduation students dissatisfied with current implementations of latent models
- © Documentation and tutorials are available on the [gensim's website](#)

LDA IN GENSIM

- © Let's first translate a set of documents (articles) into a matrix representation with a row per document and a column per feature (word or n-gram)

LDA IN GENSIM

```
from sklearn.feature_extraction.text import CountVectorizer

cv = CountVectorizer(binary      = False,
                     stop_words = "english",
                     min_df     = 3)
docs = cv.fit_transform(data.body.dropna())

# Build a mapping of numerical ID to word
id2word = dict(enumerate(cv.get_feature_names()))
```

LDA IN GENSIM

- © We want to learn which columns are correlated (i.e. likely to come from the same topic)
- © This is the word distribution
- © We can also determine what topics are in each document, the topic distribution

LDA IN GENSIM

```
from gensim.models.ldamodel import LdaModel
from gensim.matutils import Sparse2Corpus

# First we convert our word-matrix into gensim's format
corpus = Sparse2Corpus(docs, documents_columns = False)

# Then we fit an LDA model
lda_model = LdaModel(corpus = corpus, id2word = id2word, num_topics = 15)
```

LDA IN GENSIM

- © In this model, we need to explicitly specify the number of topic we want the model to uncover
- © This is a critical parameter, but there is not much guidance on how to choose it
 - © Try to use domain expertise where possible

LDA IN GENSIM

- ◉ Now we need to assess the goodness of fit for our model
- ◉ Like other unsupervised learning techniques, our validation techniques are mostly about interpretation
- ◉ Use the following questions to guide you:
 - ◉ Did we learn reasonable topics?
 - ◉ Do the words that make up a topic make sense?
 - ◉ Is this topic helpful towards our goal?

LDA IN GENSIM

- ◉ We can evaluate fit by viewing the top words in each topic
- ◉ `gensim` has a `show_topics` function for this:

```
num_topics = 25
num_words_per_topic = 5

for ti, topic in enumerate(
    lda.show_topics(num_topics = num_topics,
                    num_words_per_topic = num_words_per_topic)):
    print("Topic: %d" % ti)
    print(topic)
    print()
```

MIXTURE MODELS AND LANGUAGE PROCESSING

- ◉ Some topics will be clearer than others
- ◉ The following topics represent clear concepts:

```
0.009*cup + 0.009*recipe + 0.007*make + 0.007*food + 0.006*sugar
```

```
→ Cooking and Recipes
```

```
0.013*butter + 0.010*baking + 0.010*dough + 0.009*cup + 0.009*sugar
```

```
→ Cooking and recipes
```

```
0.013*fashion + 0.006*like + 0.006*dress + 0.005*style
```

```
→ Fashion and Style
```

ACTIVITY: KNOWLEDGE CHECK

DIRECTIONS: ANSWER THE FOLLOWING QUESTIONS

1. Demonstrate the code you used to generate the topics above
2. Hypothesise other topic interpretations



EXERCISE

INTRODUCTION

WORD2VEC

WORD2VEC

- ◉ Word2Vec is another unsupervised model for latent variable NLP
- ◉ It was originally released by Google and further refined at Stanford
- ◉ This model creates word vectors, multidimensional representations of words

```
assembly = [0.12315, 0.23425, 0.89745324, 0.235234, 0.234234, ]
```

- ◉ This is similar to having a distribution of concepts or topics that the word may come from

WORD2VEC

- ◉ If we take our usual document-word matrix and take its transpose, instead of talking about words as being features of a document, we can talk about **documents as being features of a specific word**
- ◉ In other words, how do we define or characterise a single word?
 - ◉ We can do so by defining its dictionary definition
 - ◉ Or we can enumerate all of the ways we might use it

WORD2VEC

- ◉ Given the word “Paris”, we have many contexts or uses we may find it in:

```
["_ is the capital of", "_ , France", "the capital city _", "the restaurant in _"]
```

- ◉ There are also a bunch of contexts we do not expect to find it in

```
["can I have a _", "there's too much _ on this"] # ... and millions more
```

- ◉ We could make a feature or column for each of these contexts
- ◉ We could represent “Paris” in a sparse feature with all possible contexts

WORD2VEC

- ◉ Additionally, the first few examples represent the same concept
- ◉ Paris is a city like thing, so it contains shops and restaurants
- ◉ Paris is a capital city
- ◉ We want to use dimensionality reduction to find a few concepts per word instead of all possible contexts

WORD2VEC

- ◉ With LDA, we could do this by identifying the topics a word was most likely to come from
- ◉ With Word2Vec, we will replace the overlapping contexts by some concept that represents them
- ◉ Like other techniques, our goal is to identify correlated columns and replace them with a new column that represents those replaced columns
- ◉ We can replace the columns ["_ is a city", "_ is a capital", "I flew into _ today"] by a single column IsACity

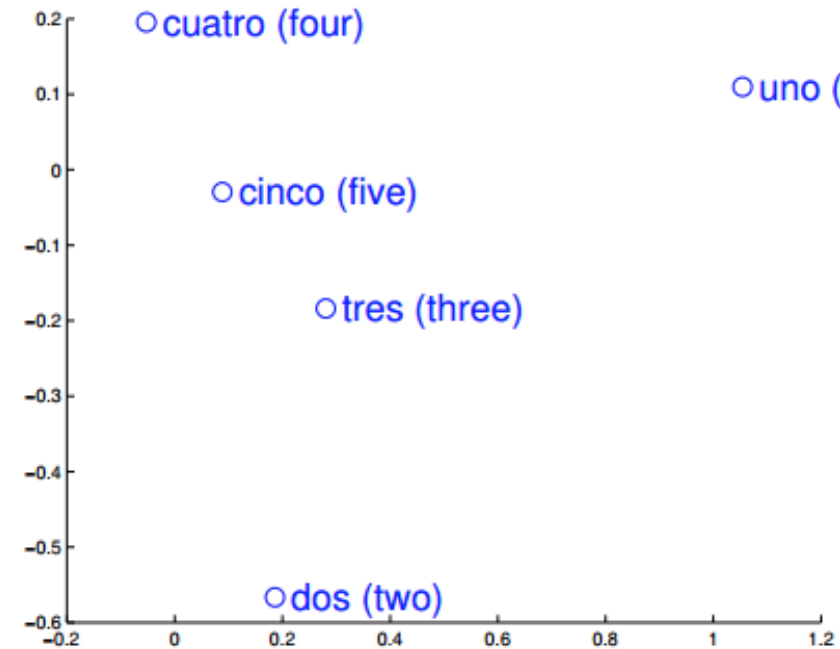


WORD2VEC

- ◉ With a trained model, Word2Vec can be used for many tasks
- ◉ A commonly used feature of Word2Vec is being able to ask what words are similar to each other
- ◉ For example, if you ask for words similar to “france”, you would get

spain	0.678515
belgium	0.665923
netherlands	0.652428
italy	0.633130
switzerland	0.622323
luxembourg	0.610033
portugal	0.577154

WORD2VEC

- © If we have data for other languages, Word2Vec could also be used for translation



ACTIVITY: KNOWLEDGE CHECK

DIRECTIONS: ANSWER THE FOLLOWING QUESTIONS

1. After reviewing the analogies, brainstorm some word vector math
 - a. For example, what do you think would happen if you subtracted the vector for “Man” from “King”?
 - b. Do you think you would get “Queen”?



EXERCISE

DEMONSTRATION

WORD2VEC IN GENSIM

WORD2VEC IN GENSIM

- © We will build a Word2Vec model using the body text of the articles available in the StumbleUpon data set

```
from gensim.models.word2vec import Word2Vec

# Setup the body text
text = data.body.dropna().map(lambda x: x.split())

from gensim.models import Word2Vec
model = Word2Vec(text, size = 100, window = 5,
                  min_count = 5, workers = 4)
```

WORD2VEC IN GENSIM

- ◉ The Word2Vec class has many arguments
 - ◉ `size` represents how many concepts or topics we should use
 - ◉ `window` represents how many words surrounding a sentence we should use as our original feature
 - ◉ `min_count` is the number of times that context or word must appear
 - ◉ `workers` is the number of CPU cores to use to speed up model training

WORD2VEC IN GENSIM

- ◉ The model has a `most_similar()` function that helps find the words most similar to the one you queried
- ◉ This will return words that are most often used in the same context

```
model.most_similar(positive = ["cookie", "brownie"])
```

- ◉ It can easily identify words related to those from this data set

INDEPENDENT PRACTICE

TWITTER LABORATORY

ACTIVITY: TWITTER LABORATORY



EXERCISE

DIRECTIONS: (45 MINUTES)

1. In this exercise, we will compare some of the classical NLP tools from the last class with these more modern latent variable techniques
2. We will do this by comparing information extraction on Twitter using two different methods
 - a. NOTE: There is a pre-existing file of captured tweets you can use
 - b. It is located in the class repository for lesson-14
 - c. However, you can also collect your own tweets following the instructions in `twitter-instructions.md`
3. Check
 - a. `~/lessons/lesson-14/code/starter/*`

ACTIVITY: TWITTER LABORATORY



EXERCISE

DIRECTIONS: (45 MINUTES)

1. Use `spaCy` to write a function to filter tweets down to those where Google is announcing a product. How might we do this? One way might be to identify verbs, where "Google" is the noun and there is some action like "announcing"
 - a. Write a function that can take a sentence parsed by `spaCy` and identify if it mentions a company named "Google". Remember, `spaCy` can find entities and code them as ORG if they are a company.
 - b. BONUS: Make this function work for any company
 - c. Write a function that can take a sentence parsed by `spaCy` and return the verbs of the sentence (preferably lemmatised)
 - d. For each tweet, parse it using `spaCy` and print it out if the tweet has "release" or "announce" as a verb
 - e. Write a function that identifies countries. HINT: the entity label for countries is GPE (or "GeoPolitical Entity")
 - f. Re-run (d) to find country tweets that discuss "Iran" announcing or releasing

ACTIVITY: TWITTER LABORATORY



EXERCISE

DIRECTIONS: (45 MINUTES)

1. Build a [word2vec](#) model of the tweets we have collected using [gensim](#)
 - a. First take the collection of tweets and tokenise them using [spaCy](#)
 - i. Think about how this should be done
 - ii. Should you only use upper-case or lower-case?
 - iii. Should you remove punctuations or symbols?
 - b. Build a [word2vec](#) model
 - i. Test the window size as well - this is how many surrounding words need to be used to model a word. What do you think is appropriate for Twitter?
 - c. Test your [word2vec](#) model with a few similarity functions
 - i. Find words similar to “Syria”
 - ii. Find words similar to “war”
 - iii. Find words similar to “Iran”
 - iv. Find words similar to “Verizon”
 - d. Adjust the choices in (b) and (c) as necessary

ACTIVITY: TWITTER LABORATORY

DIRECTIONS: (45 MINUTES)

1. Filter tweets to those that mention “Iran” or similar entities and “war” or similar entities
 - a. Do this using just [spaCy](#)
 - b. Do this using [word2vec](#) similarity scores



EXERCISE

CONCLUSION

TOPIC REVIEW

TOPIC REVIEW

- ◉ Latent variable models attempt to uncover structure from text
- ◉ Dimensionality reduction is focused on replacing correlated columns
- ◉ Topic modelling (or LDA) uncovers the topics that are most common to each document and then the words most common to those topics
- ◉ Word2Vec builds a representation of a word from the way it was used originally
- ◉ Both techniques avoid learning grammar rules and instead rely on large data sets. They learn based on how the words are used, making them very flexible

DATA SCIENCE

BEFORE NEXT CLASS

BEFORE NEXT CLASS

DUE DATE

- ◉ Project:
 - ◉ Final Project, part 3

LATENT VARIABLES AND NLP

Q & A

EXIT TICKETS

DON'T FORGET TO FILL OUT YOUR EXIT TICKET

[Exit Ticket Link](#)

What's the lesson number?	14
What was the topic of the lesson?	Latent Variables and NLP

CREDITS AND REFERENCES

CREDITS AND REFERENCES

- ◉ Natural Language Processing
 - ◉ [Five Takeaways on the State of Natural Language Processing](#)
- ◉ LDA (Latent Dirichlet Allocation)
 - ◉ [Introduction to LDA](#)
 - ◉ [Gensim LDA documentation](#)
- ◉ Word2Vec
 - ◉ [Gensim Word2Vec documentation](#)
 - ◉ [How Google Converted Language Translation Into a Problem of Vector Space Mathematics](#)
 - ◉ [A word is worth a thousand vectors](#)
 - ◉ [Word Embeddings For Fashion](#)
 - ◉ [Building the Next New York Times Recommendation Engine](#)